

# AWS PostgreSQL RDS

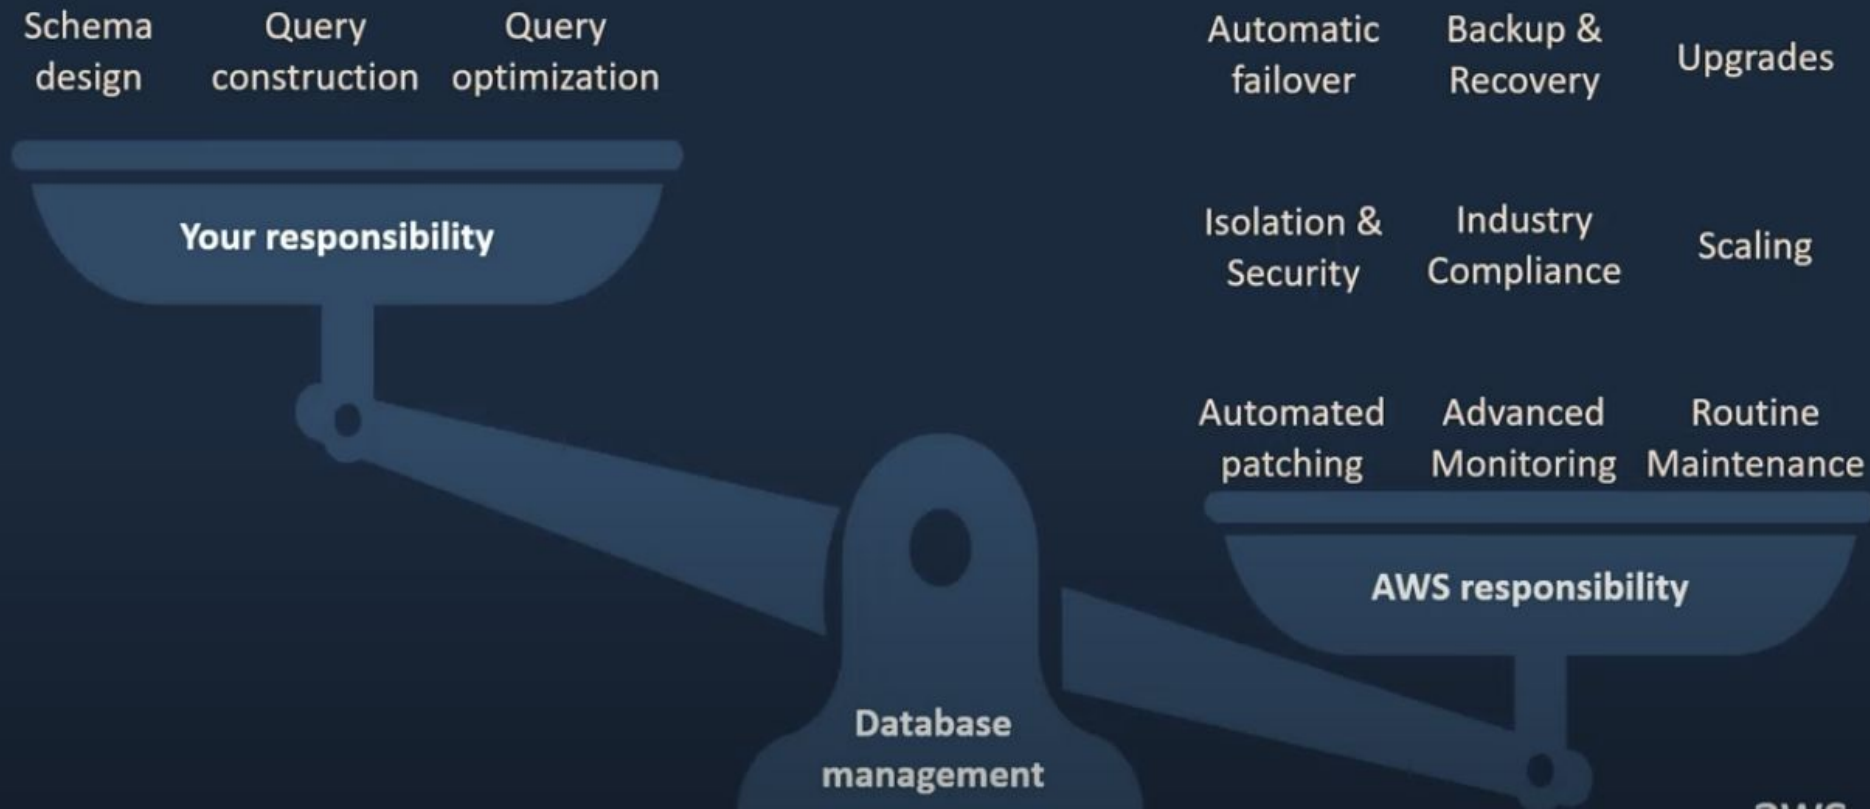
Partition

# PostgreSQL – Open Source Database

- PostgreSQL is a powerful, open source object relational database system
- Origins of PostgreSQL date back to 1986 as part of the POSTGRES project at the University of California at Berkeley
- More than 30 years of active development on the core platform
- Rich features and extensions, reliability and standards compliance, open source license

# Accelerate path to innovation with managed databases

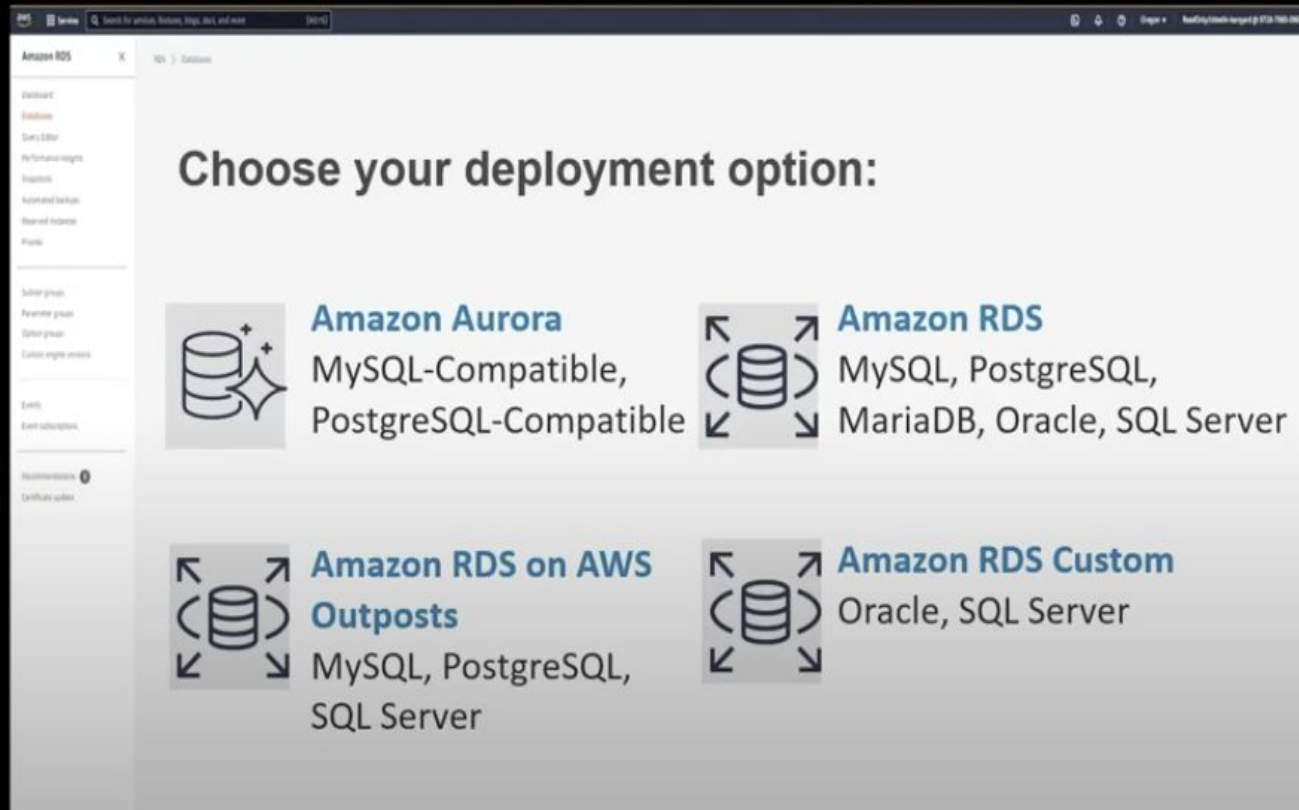
Spend time innovating and building apps, not managing infrastructure



# Deployment options with Amazon RDS

Get started with your choice of 7 popular relational databases across 4 deployment options

Unlock better  
performance,  
availability,  
scalability, and  
security



# Partition

## Challenges of hyper-scale customers

***New growth pattern:*** sudden, unanticipated, and sustained

### Challenges

- Rapid growth of data size creating large size table(s)
- Increase on-disk footprint approach infrastructure limits
- Increase query response time
- Operational challenges for vacuum, index management, upgrades, and etc.
- Increase recovery time size of large table



# The solution – Table partitioning

*Divide a large table into multiple small segments, called partitions*

## Benefits

- **Availability:** partition level backup / restore
- **Manageability:**
  - Partition-based maintenance (index, vacuum)
  - Easy add / remove partitions via partitioning “sliding window” for data lifecycle management
- **Performance:**
  - Allow query execution to target a partition / subset of partitions (partition pruning)
  - Partition wise join

# Table partitioning in PostgreSQL

## <= PG v9.6

Table  
Inheritance

CHECK  
Constraints

Trigger for  
routing

## PG v10

Declarative  
Table  
Partitioning

## PG v11

Hash  
Partitioning

Indexes on  
partitioned  
tables

FOR EACH ROW  
trigger

Partition wise  
join

Partition wise  
aggregates

## PG v12

Partition Pruning  
enhancements

Foreign keys to  
reference  
partitioned tables

Faster Copy for  
partitioned tables

Non-blocking  
ATTACH Partition

## PG v13

Partition wise  
join  
improvements

Partition  
Pruning  
improvements

Before row-  
level triggers

Support logical  
replication

## PG v14

Detach partition  
concurrently

Improve update /  
delete  
performance

# Declarative table partitioning

- A partitioned table is a virtual table without any physical storage
- Physical data is stored in partition leaf tables
- Data is split into partitions based on the partition strategy defined
  - Partitioning method
  - Partitioning key
- The subset of data in a partition is defined by Partition Bounds
- A partition may contain sub-partitions
- A default partition stores data that does not belong to any partitions



# Partitioning strategy – Partitioning method

**Decision point:** partitioning method

## Partitioning methods

- **Range:** Data is placed in partitions based on a range of values
  - Lower end: inclusive, upper end: exclusive
- **List:** Data is placed in partitions based on a list of discrete values
- **Hash (PG v11+):** Data is placed in partitions based on a hash algorithm applied to a key
- **Composite:** Combining one or more of partitioning types, e.g. list + range, hash + list, and etc...

# Partitioning strategy – Partition key

**Decision point:** partition key

## Know your data & workload access patterns

- The key must be made up of columns and/or immutable expression
- The key value must be a literal
- The partition key can include up to 32 columns or expressions
  - List partitioning is limited to a single column

# Time-series data

- Data is treated as immutable append-only log, table grows continuously overtime
- Old rows may be deleted if purging is needed, but they are not changed
- Typical access to data is specific for a time range
- User cases: Internet of things device streaming data, transaction history table in an OLTP database

## *Relational database with partitioning vs. purpose-built database*

### Amazon RDS for PostgreSQL

- Need to perform join with other relational data
- Data can be injected in batches

### Purpose-built Amazon Timestream

- Time-series data to be used in standalone application
- Data is loaded directly from many locations or sensors

# Example

**Scenario:** online delivery app experiencing sudden, rapid growth of demands, 10 TB transaction history table growing at 2 TB a week.

- Need to control size of the table
- Vacuum process is not able to keep up w/ the change rate
- Increasing query response time
- Customers access to orders in the last 90 days only, orders older than 90 days can be archived

```
CREATE TABLE transaction_history (  
    time                timestampz,  
    transaction_id      bigint,  
    total               money,  
    status              text,  
    transaction_details jsonb  
);
```

# The solution – Range partitioning

## Decision on partitioning strategy

- Partition method: Range
- Partition key: time
- Partition boundary: day (transaction history older 90 days will be archived)

```
CREATE TABLE transaction_history (  
    time                timestamptz,  
    transaction_id      bigint,  
    total               money,  
    status              text,  
    transaction_details jsonb  
) PARTITION BY RANGE(time);
```



# Manage partitions (native PostgreSQL)

Native PostgreSQL does not provide ways to manage partition automatically

- Create partitions manually
- Use cron scripts

```
CREATE TABLE transaction_history_p2022_01_01
PARTITION OF transaction_history
FOR VALUES FROM ('2022-01-01 00:00:000') TO ('2022-01-02 00:00:000');
```

```
CREATE TABLE transaction_history_p2022_01_01
PARTITION OF transaction_history
FOR VALUES FROM ('2022-01-02 00:00:000') TO ('2022-01-03 00:00:000');
```

```
# repeat to create partitioning for 90 days ...
```

# Automate partition management

## pg\_partman

- An extension to create and manage both time-based and serial-based table partition sets
- RDS for PostgreSQL support starts w/ PG v12.5

## pg\_cron

- An extension that allows you to use cron syntax to schedule PostgreSQL commands directly within your database
- RDS for PostgreSQL support starts w/ PG v12.5

# pg\_partman

## Automate operations for

- Create partition parent table
- Continuous creation of partition child tables, attach to partitioned table
- Attach new partitions
- Data lifecycle management (detach/drop aged partitions)

*Enable pg\_partman separately for each database:*

```
CREATE SCHEMA partman;
```

```
CREATE EXTENSION pg_partman WITH SCHEMA partman;
```

# Automate partition creation

## Policy-based automatic partition management

- Create\_parent() function call
- Define policy for creating new partitions continuously

```
SELECT partman.create_parent(  
    'public.transaction_history',  
    'time',  
    'native',  
    'daily',  
    p_start_partition := (now() - interval '90 days')::date::text );
```

# Ongoing partition maintenance

Manage partitions in rolling window: create new / remove old partitions

- Set retention policy with `part_config`
- Define actions on “aged” partitions
- Call `run_maintenance_proc()` to trigger maintenance

```
UPDATE partman.part_config
  SET retention = '90 days',
      retention_keep_table=true
 WHERE parent_table = 'public.transaction_history';
```



# Automate maintenance with pg\_cron

- To enable pg\_cron on RDS for PostgreSQL
  - Add pg\_cron to shared\_preload\_libraries parameter in the DB instance's parameter group
  - Create extension
- Set scheduled maintenance with calling run\_maintenance\_proc()

```
CREATE EXTENSION pg_cron;  
SELECT cron.schedule(  
    'transaction_history maintenance',  
    '0 22 * * *',  
    $$CALL partman.run_maintenance_proc()$$  
);
```

# Results – Reduce query response time

*Results of query pruning – a query optimization technique that improves performance for declaratively partitioned tables*

If the query is constructed properly,  
PostgreSQL can eliminate whole partitions  
from being scanned

```
SELECT sum(total)
FROM transaction_history
WHERE time
BETWEEN '2022-09-26'
AND '2022-09-27';
```

2022-09-24

2022-09-25

2022-09-26

2022-09-27

2022-09-28

2022-09-29

2022-09-30

2022-10-01

# Results – Faster bulk loads and deletes

## Bulk loads

- Create stand-alone table without indexes or constraints
- Perform bulk load
- Create indexes or constraints
- ALTER TABLE ATTACH PARTITION ... to add to partitioned table
- **Benefits:** reduce locking and impacts to parent table during bulk loads

## Bulk deletes

- ALTER TABLE DETACH PARTITION, then Drop TABLE
- **Benefits:** avoid the VACUUM overhead caused by a bulk DELETE

# Results – Data lifecycle management

## Migrate less frequent used data to cheaper storage

- Rolling windows of data allows easy archival of cold data
- PostgreSQL v14+ supports DETACH PARTITION CONCURRENTLY
- Does not create bloat with deleting expired rows

2022-09-23

2022-09-24

2022-09-25

2022-09-26

2022-09-27

2022-09-28

2022-09-29

2022-09-30

2022-10-01

# How effective is the solution?

- ✓ Manage size of the table
  - Data divided into child tables with less than 3 GB per partition
  - Partitioned table stores most recent 90 days of data
- ✓ Reduce query response time, enhance overall performance
  - Allow query execution to target a subset of partitions with query pruning
  - Partition wised join
  - Partition-based index maintenance, vacuum, backup / restore
- ✓ Enable easy archival of historical data
  - Automatic archival of data older than 90 days
  - Concurrent detach of partitions with aged data without creating bloat



# When to use partitioning?

*For very large tables (e.g. > 100 million rows)*

- Size of table unable to fit in memory
- For time-based or series-based data
  - Continuously streaming of data in append only fashion
  - Historical data can be archived or discarded
- Data access pattern
  - Querying data by time range, or discrete values
  - Data injection in batches

# What partitioning is not?

## *Not a magic bullet*

- Not for every table, partition has performance and operational overhead
- No substitute for good schema design or operation best practices
- Not substitute for query tuning
- Not sharding



# Best practices – Partition key

## *know your data & workload access patterns*

- Column(s) which most commonly appear in WHERE clauses for effective query pruning
- Column(s) to support rolling windows of data for archiving
- Column(s) that do not change often, to avoid moving rows among partitions

# Best practices – Number of partitions

- Too few: individual partitions is too big
- Too many: increase query planning time
- Use sub-partitioning to divide data into multi-levels for very large data sets



# Best practices

- Start partitioning early before size of table become unmanageable
- Get on the latest major and minor version releases for partitioning enhancements
- Use `pg_partman` and `pg_cron` extension to automate partition management



# Limitations

- Primary keys must contain the partition key
- Unique Indexes: cannot create unique index across partitions.
- When an UPDATE causes a row to move from one partition to another, there is a chance that another concurrent UPDATE or DELETE will get a serialization failure error

# Summary

- Partitioning benefits
  - Improved query performance (query pruning, partition-wise join)
  - Faster bulk load and bulk delete operations
  - Data lifecycle management
- Knowing your data access pattern helps to create effective partitioning strategy
- Implementing partitioning is transparent to clients
- Partitioning adds additional operational overhead and maintenance burden, evaluate if partitioning is the right solution for you